



Building Cloud-Native Applications

Yulia Tenincheva

Senior Cloud Engineer, MentorMate

What is expected of you?

- To know how to develop efficient and secure Node.js applications that solve the business need
- To know how to use AWS SDK to use some of the more popular cloud services
- To adopt a [Cloud Mindset](#)
- To understand the Serverless architecture
- To continue learning and growing your skillset
- And if you get AWS Certified one day - I will be extremely proud of you :D

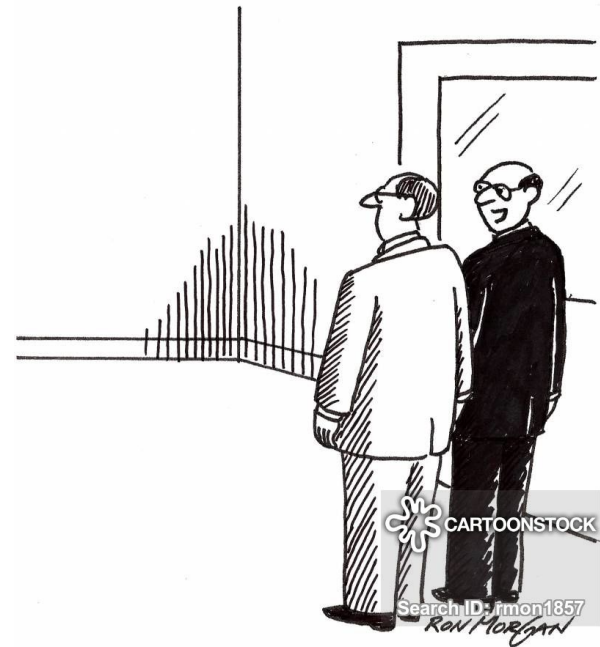


Cloud Native

Cloud native is an **approach** to building and running applications that fully exploit the **advantages** of the **cloud computing** model.

Cloud-Native Apps Are Different

Services	<ul style="list-style-type: none">• All functionality is published and consumed via web services
Handling Failures	<ul style="list-style-type: none">• Every Integration point will eventually fail one time or another• Be prepared to handle all kind of failures
Horizontal Scalability	<ul style="list-style-type: none">• Design for Scale Out
Asynchronous Processing	<ul style="list-style-type: none">• Break down the task, process requests asynchronously• Use queues to decouple functionality• Eventual consistency model
Stateless Model	<ul style="list-style-type: none">• Build stateless services that can be scaled out and load balanced
Minimize Human Intervention	<ul style="list-style-type: none">• Go DevOps/NoOps

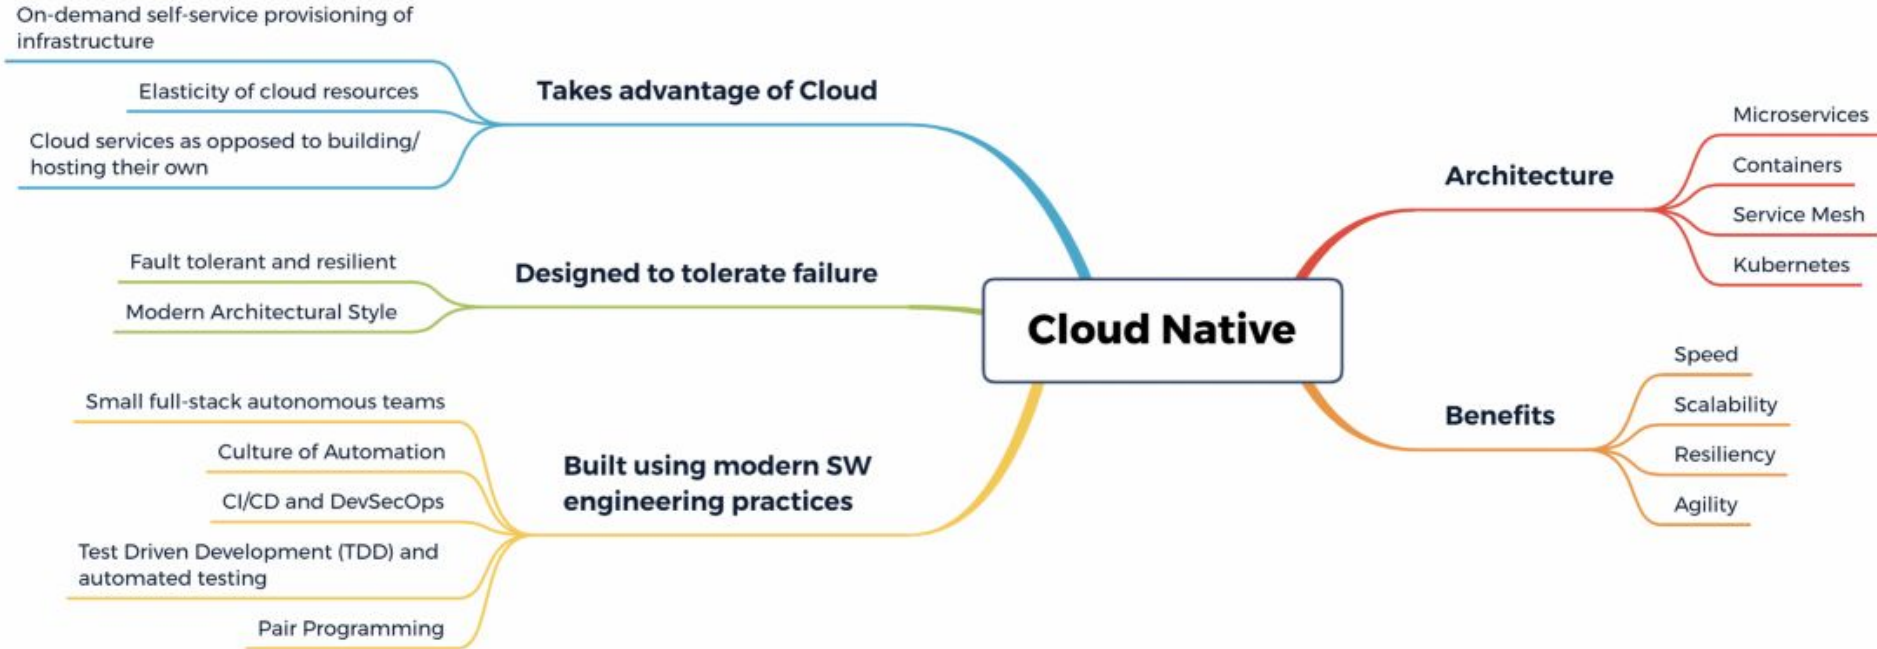


"Everything is in the Cloud."

“Everything fails all the time”

~ Werner Vogels, Vice President & CTO at Amazon.com

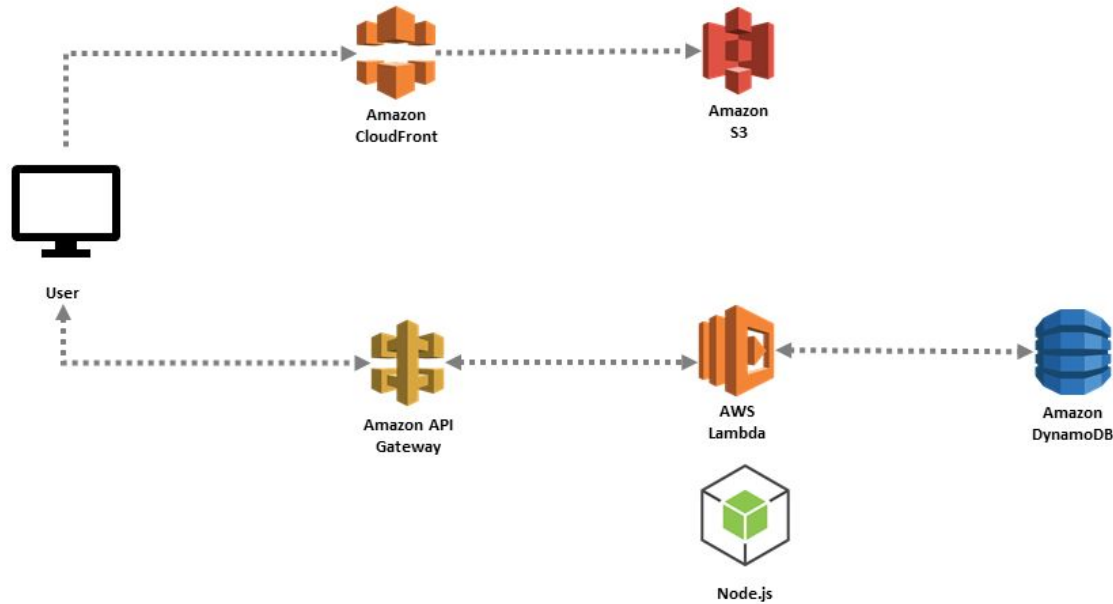
Cloud-Native Design Patterns / Mindset



Cloud-Native Design Patterns / Mindset

- Cloud-Native apps need rapid scale
- At scale failures are inevitable
- Let apps handle own resiliency
- Build stateless services
- Scale out, not scale up
- Treat infrastructure differently
- Create immutable infrastructures
- Adopt Microservices Architecture

Base AWS Services for Serverless



S3

S3 (Simple Storage Service)

Amazon Simple Storage Service (Amazon S3) is an **object storage service** that offers industry-leading scalability, data availability, security, and performance.

Customers can use it **to store and protect any amount of data** for a range of use cases, such as websites, mobile applications, backup and restore, archive, enterprise applications, IoT devices, and big data analytics.

Amazon S3 is designed for **99.999999999%** (11 9's) of **durability**, and stores data for millions of applications for companies all around the world.



Concepts of S3

- Bucket

- Container for objects stored in S3
- Unlimited size
- Organize the namespace at the highest level
- Internet accessible storage via HTTP/s
- Global unique name
- Limit of 100 buckets per account

- Key


- Name of the object
- Unique identifier for an object within a bucket
- Use the object key to retrieve the object

- Object

- Similar to files
- No hierarchy
- Objects are immutable
- Size up to 5 TB
- Uniquely identified within a bucket by a key(name) and a version ID

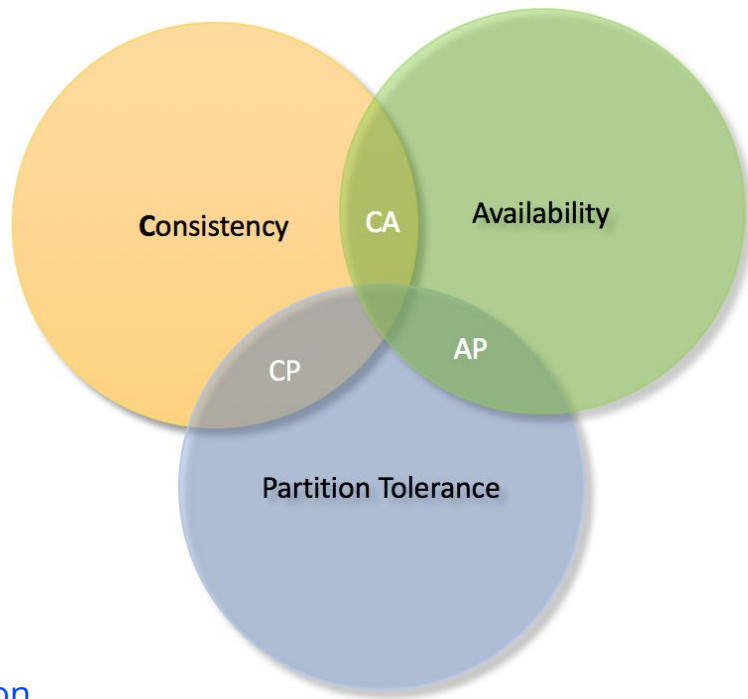
How to Access S3

- AWS CLI
 - `aws s3 help`
- REST API
 - Get / Put Object
- [AWS SDK](#)
 - `getObject`, `getSignedUrl`, `headObject`, ...
 - `putObject`, `upload`, `uploadPart`, ...
- Web Console
 - Like a monkey



CAP Theorem

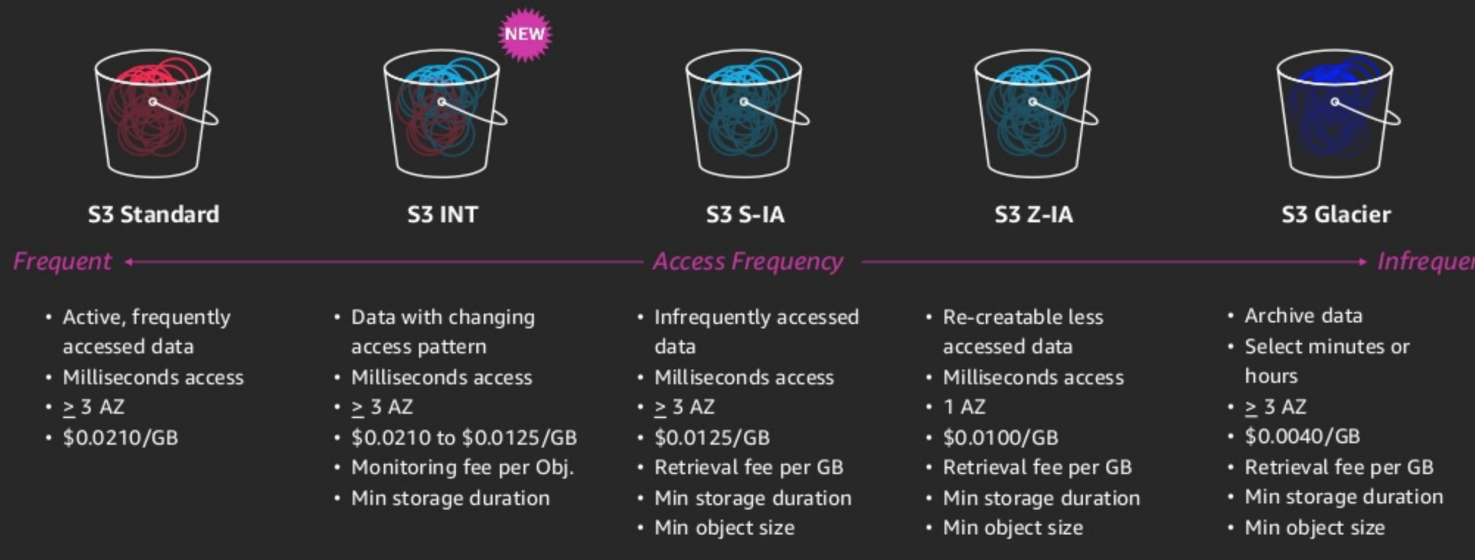
- Amazon S3 (AP)
 - **High availability**
 - **High durability** (99.999999999%)
 - **Eventual consistency** for overwrite and deletes
 - **Read-after-write consistency** for new uploads



AWS S3 Consistency model in [details](#) & [quick explanation](#)

S3 Storage Classes

Your choice of Amazon S3 storage classes



S3 Features

- Versioning
- Static Website Hosting
- Cross-Region Replication
- Lifecycle Management
- Encryption
- Security - ACLs and Bucket Policies




CLOUD STORAGE CHEAT SHEET



A CLOUD GURU

FILE STORAGE

- NETWORK FILE SYSTEMS
- CAN BE SHARED AMONG MULTIPLE COMPUTE INSTANCES



EXAMPLE SERVICES:

AMAZON EFS (LINUX)

AMAZON FSx (WINDOWS)

GOOGLE CLOUD FILESTORE

AZURE FILE STORAGE



COMMON "VOLUME" MOUNTED AS VIRTUAL DISK

HIGH-PERFORMANCE DATA ACCESS FOR A SINGLE SERVER

BLOCK STORAGE

OBJECT STORAGE

COMMON PROTOCOL: API

NEAR-INFINITE SCALE
HIGHLY AVAILABLE
EXTREMELY DURABLE

HOT — FAST ACCESS, MORE EXPENSIVE
or COOL — SLOWER + CHEAPER — GREAT FOR BACKUPS

AMAZON S3

MANY STORAGE TIERS
HOT → COOL

AZURE BLOB STORAGE
HOT OR COOL

GOOGLE CLOUD STORAGE
HOT!
COOL!

GCS NEARLINE

@forrestbrazeal

COLD STORAGE

- ARCHIVAL DATA
- TAPE BACKUPS
- SLOW, BUT CHEAP

AMAZON S3 GLACIER

GOOGLE CLOUD STORAGE COLDLINE

AZURE ARCHIVE BLOBS



MENTORMATE™
DevCamp

Static **WebSite** Hosting in S3 - Demo

AWS CloudFront

What is a CDN and why use one?

- Global Content Delivery Network with massive capacity and scale
- Optimized for Performance and scale
- Built-in Security features
- Robust Real Time reporting
- Static and Dynamic object, video delivery

CloudFront: An Integral Part of AWS

Mobile Application Delivery

CloudFront, WAF, Route 53

Static and Dynamic Object Origin

CloudFront, WAF, Route 53, Elastic Transcoder

Web and Application Server Origin

CloudFront, WAF, Route 53, Elemental / Elastic Transcoder

Enterprise Applications

CloudFront, WAF, Route 53



AWS
Mobile Hub



Amazon
Cognito



Amazon API
Gateway



Amazon
S3



Amazon EFS



Elastic Load
Balancing



Amazon EC2



AWS
Lambda




Amazon
WorkMail



CloudFront Service Components

- Distributions
- Origins
- Behaviors
- Restrictions, Error pages, Tags
- AWS WAF Web ACLs
- Edge Locations
- Price Classes



AWS Lambda

AWS Lambda release history



*As of October 2018, does not include region launches

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved



Serverless applications

Event source



Changes in
data state



Requests to
endpoints



Changes in
resource state



Function



Services (anything)



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved

Common Lambda use cases



Web Applications

- Static websites
- Complex web apps
- Packages for Flask and Express



Backends

- Apps & services
- Mobile
- IoT



Data Processing

- Real time
- MapReduce
- Batch




Chatbots

- Powering chatbot logic



Amazon Alexa

- Powering voice-enabled apps
- Alexa Skills Kit



IT Automation

- Policy engines
- Extending AWS services
- Infrastructure management

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved



Use cases

- When to use AWS Lambda
 - Working on aws resources directly
 - Building small one off applications
 - When you need to process something after an event that occurs on AWS
- When **NOT** to use AWS Lambda
 - When operating system resources are needed
 - When native libraries are needed (e.g ffmpeg)
 - When the application is large and complex
 - When you need lots of memory
 - When you need lots of execution time



Use cases

- Questionable
 - When your function is very complex
 - When your function takes an unknown amount of time to finish
- When in doubt:
 - Get familiar with the things you will be using
 - Build a PoC (Proof of concept) in a small and controlled way



Anatomy of a Lambda function

Handler() function

Function to be executed upon invocation

Event object

Data sent during Lambda function Invocation

Context object

Methods available to interact with runtime information (request ID, log group, more)

```
public String handleRequest(Book book, Context context) {  
    saveBook(book);  
  
    return book.getName() + " saved!";  
}
```

© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved




Lambda execution model

Synchronous (push)



Asynchronous (event)



Poll-based



© 2019, Amazon Web Services, Inc. or its Affiliates. All rights reserved



Source of these Slides and full lecture:

https://www.youtube.com/watch?v=EBSdyoO3goc&ab_channel=AmazonWebServices


AWS Lambda Best Practices

- Limit your function size
- Node.js - remember executions is asynchronous
- Don't assume function container reuse - but take advantage of it when it does occur
- Don't forget about disk (500Mb /tmp directory provided to each function)
- Create custom metrics (operations-centric and business-centric)



One more thing

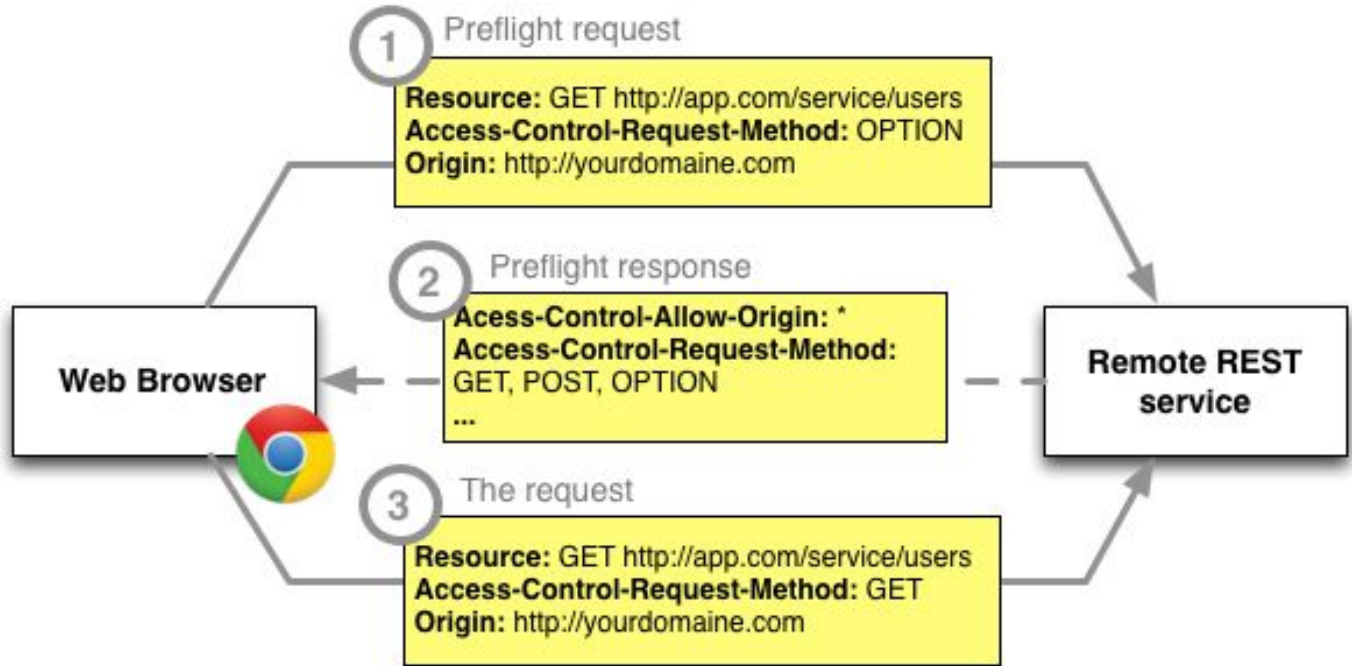
- AWS Lambda Pricing
 - The AWS Lambda **free usage** tier includes **1M free requests per month** and **400,000 GB-seconds** of compute time per month.
 - In the AWS Lambda resource model, you choose the amount of memory you want for your function, and are allocated proportional CPU power and other resources. An increase in memory size triggers an equivalent increase in CPU available to your function.



Wire S3 with AWS Lambda via API Gateway (Demo)

CORS

CORS



- Everything you would like to know about Cross Origin Resource Sharing - [here](#)

Homework


Homework

1. Host a static web site in your bucket (it can be a plain “hello world”)
2. Create a web server with a single endpoint POST /upload. Use [Postman](#) as a client to [invoke](#) and *test* your endpoint. In your server you should write a handler that processes the `multipart/form-data` request and re-upload that file to S3.

Tricky parts & considerations:

- Keep your access/secret keys... secret!
Do not expose them in any way!
- Make sure that the permissions of the file does not allow public access to it.
- Test your API with bigger files. Is it working efficiently? Consider implementing [multipart](#) uploads.
- Can you apply streams and pipes here? If so, do it.

Final requirement: When the file upload is ready and successful, return to the client a *signed url* of the uploaded file that. That URL should expire after 15 minutes.



Homework

3. (Optional) Following [this](#) example, create an AWS Lambda function that resizes images. Make it so that function is triggered automatically whenever a new file is uploaded to your bucket source “*folder*”.

Name your AWS Lambda function uniquely, so it is easy to identify you as the author. Configure IAM permissions of that Lambda in such a way, so it can access your bucket only.

Important: Write the result to another “*folder*” in S3 to avoid circular invocations (infinite loop)!!

If you do this exercise, you will learn how to import NPM modules to your Node.js Lambda functions, how to use the AWS CLI to deploy your code, how to configure fine-grained access and you will also learn how to configure asynchronous events.

4. (NOT a homework!! This is for Cloud enthusiasts only) - Wild Rydes [hands-on workshop](#)* for serverless architecture.

*You can use S3 + CloudFront, instead of AWS Amplify.

Please, destroy the created resources when you are done with the workshop.

