



Serverless

Yulia Tenincheva

Senior Cloud Engineer, MentorMate

Serverless Architecture

How should my app
withstand a server failing?

How can I tell if a
server has been
compromised?

How can I increase
utilization of my servers?

Which OS should my
servers run?

How much remaining
capacity do my servers have?

How should I implement dynamic
configuration changes on my servers

When should I decide to
scale up my servers?

What size servers are
right for my budget?

How will I keep my server
OS patched?

How can I control
access from my servers?

Which packages should
be baked into my server images?

Servers

(AAHHHHHHHHH!!)

How will new code be
deployed to my
servers?

How will the application
handle server hardware failure?

How many users create
too much load for my servers?

What size server is
right for my performance?

Which users should have
access to my servers?

Should I tune OS settings
to optimize my application?


When should I decide to
scale out my servers?

How many servers
should I budget for?



Architect to be Serverless

- Fully Managed
 - No provisioning
 - Zero administration
 - High availability
- Developer Productivity
 - Focus on the code that matters
 - Innovate rapidly
 - Reduce time to market
- Continuous Scaling
 - Automatically
 - Scale up and scale down



Many Serverless Options on AWS



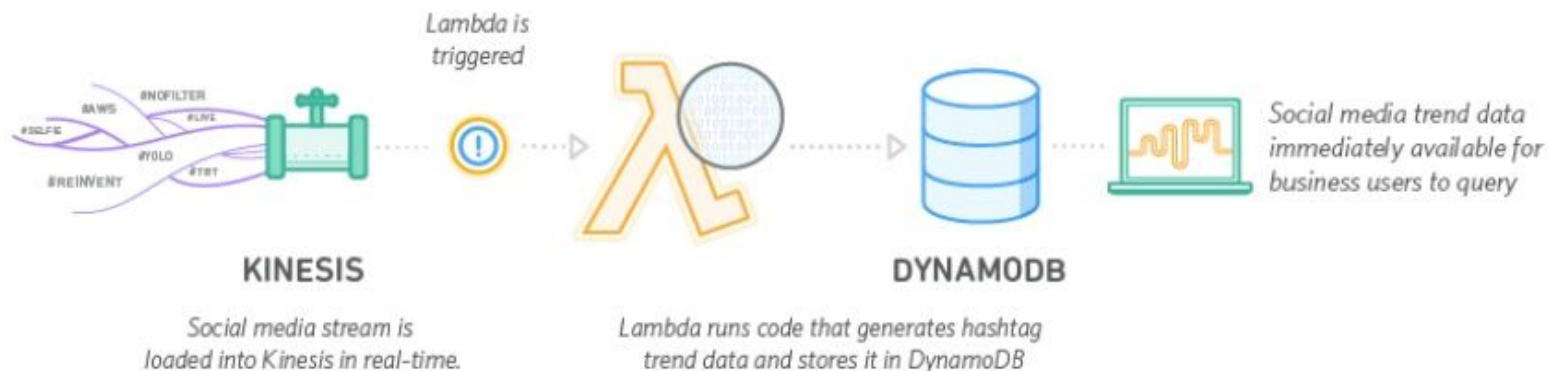
Example of Serverless Architecture

Example: Weather Application

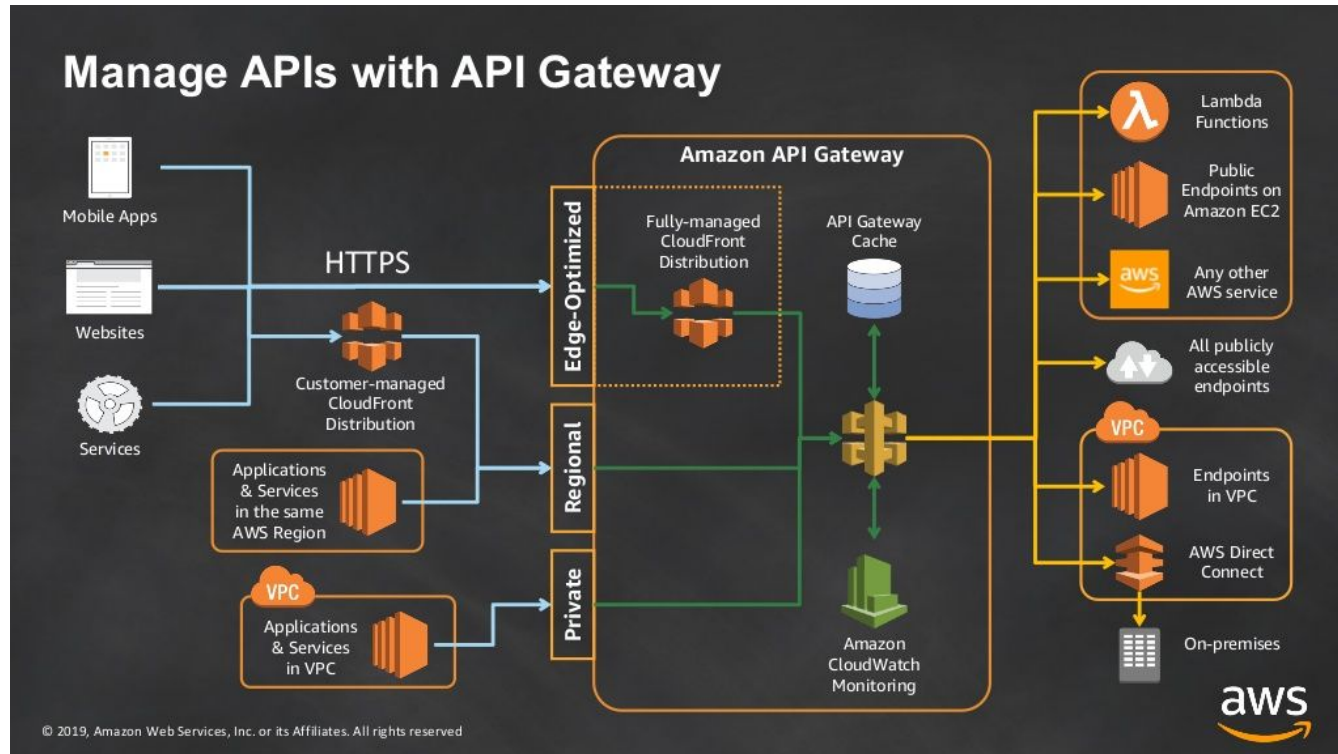


Example of Serverless Architecture

Example: Analysis of Streaming Social Media Data



Example of Serverless Architecture



Serverless Building Blocks

Database



Amazon
DynamoDB

Gateways



Amazon
API Gateway

Security



AWS
IAM



AWS
KMS

Messaging and Queues



Amazon
SQS



Amazon
SNS

Compute



AWS Lambda

Storage



Amazon S3

Network



Amazon
VPC



Amazon
Route 53



Elastic Load
Balancing

Content Delivery



Amazon
CloudFront

Streaming Analytics



Amazon Kinesis

User Management



Amazon Cognito

Internet of Things



AWS IoT

Monitoring & Logging



Amazon
CloudWatch

Machine Learning

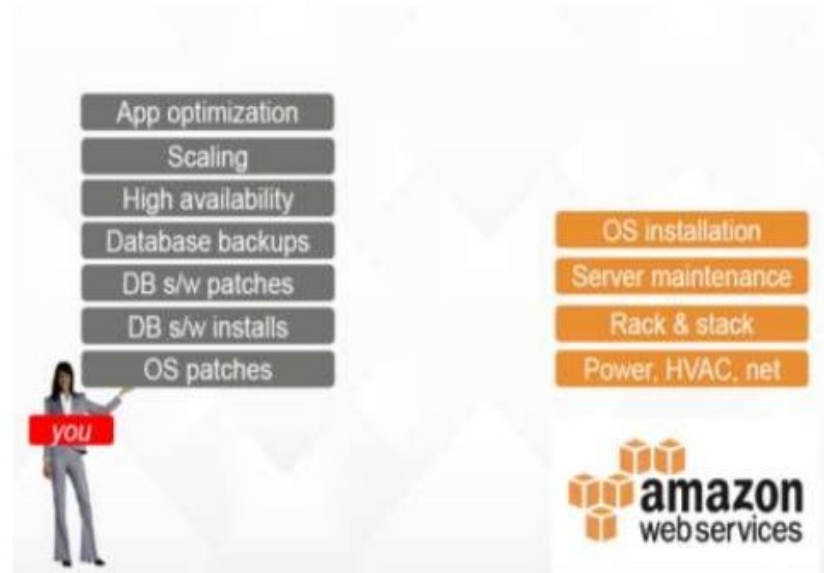


Amazon
Machine Learning

Serverless Database ?









DB hosted on-premises



DB hosted on Amazon EC2

DynamoDB Benefits

-  Fully managed
-  Fast, consistent performance
-  Highly scalable
-  Flexible
-  Event-driven programming
-  Fine-grained access control

Duolingo Scales to Store Over 31 Billion Items Using DynamoDB

“

Using AWS, we can handle traffic spikes that expand up to seven times the amount of normal traffic.

Severin Hacker
CTO, Duolingo

duolingo.

”

Duolingo is a free language learning service where users help translate the web and rate translations.

- Duolingo stores data about each user to be able to generate personalized lessons.
- The MySQL database couldn't keep up with Duolingo's rate of growth
- By using the scalable database service, data store capacity increased from 100 million to more than four billion items
- Duolingo has the capacity to scale to support over 8 million active users

Source: [This case study](#)

What is DynamoDB?

- NoSQL database **tables** as a service
- Store as many **items** as you want
- Items may have different **attributes**
- **Low-latency** queries
- Scalable read/write **throughput**

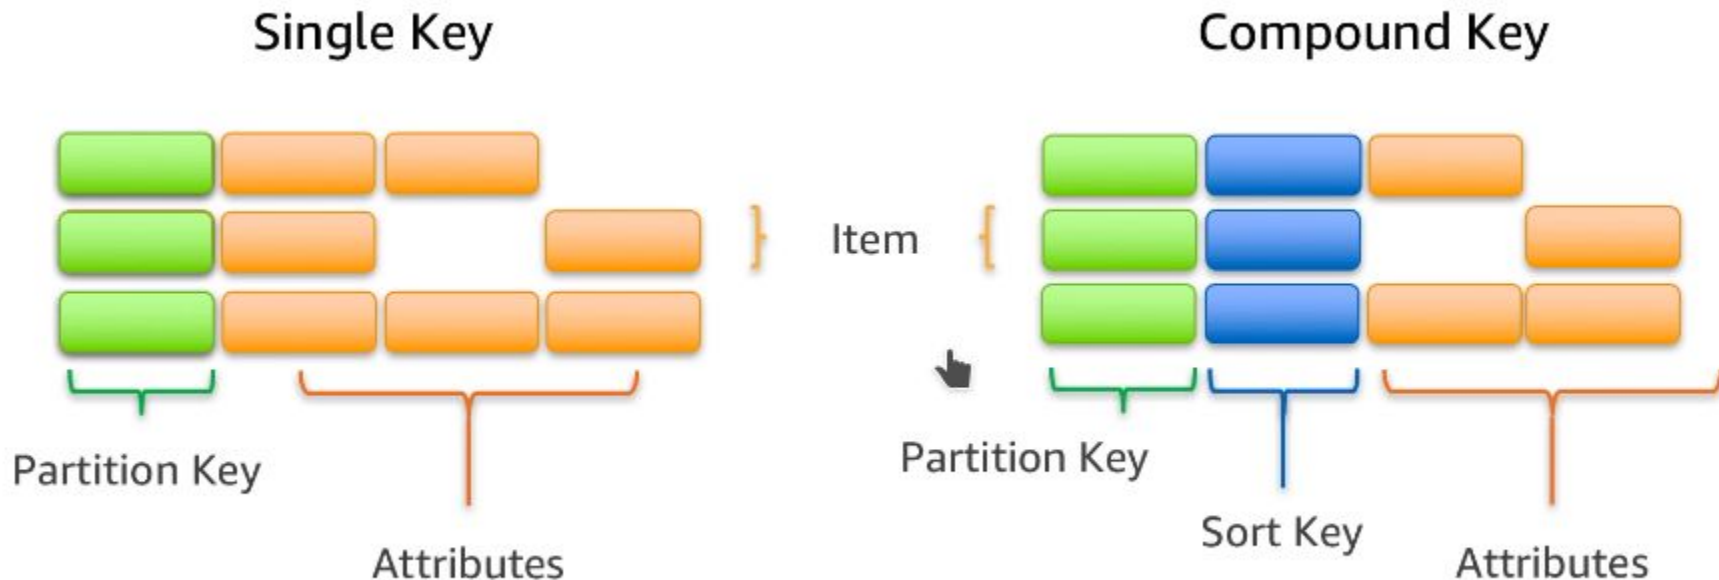
People

```
{
  "PersonID": 101,
  "LastName": "Smith",
  "FirstName": "Fred",
  "Phone": "555-4321"
}

{
  "PersonID": 102,
  "LastName": "Jones",
  "FirstName": "Mary",
  "Address": {
    "Street": "123 Main",
    "City": "Anytown",
    "State": "OH",
    "ZIPCode": 12345
  }
}

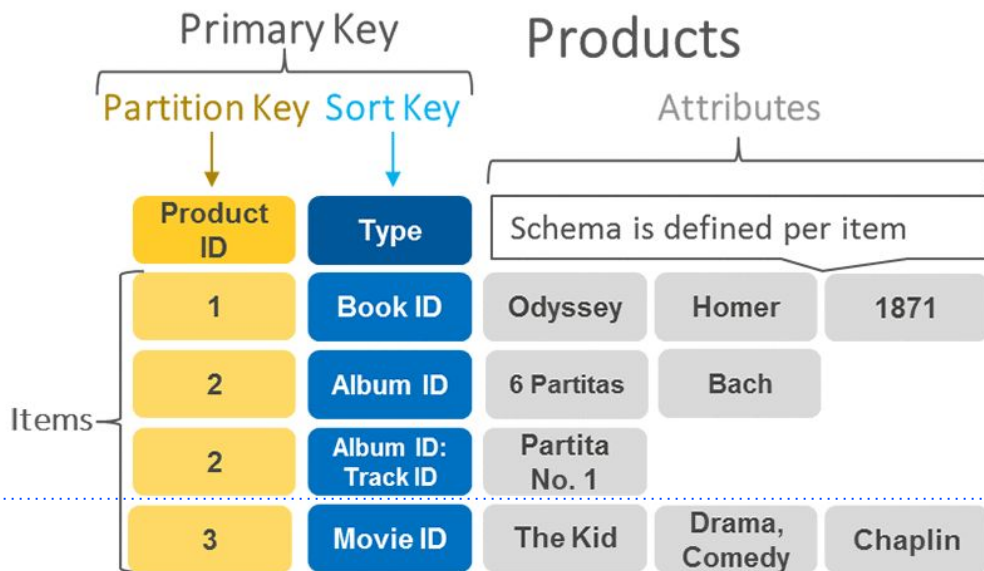
{
  "PersonID": 103,
  "LastName": "Stephens",
  "FirstName": "Howard",
  "Address": {
    "Street": "123 Main",
    "City": "London",
    "PostalCode": "ER3 5K8"
  },
  "FavoriteColor": "Blue"
}
```

Items in a table must have a key



Partition key

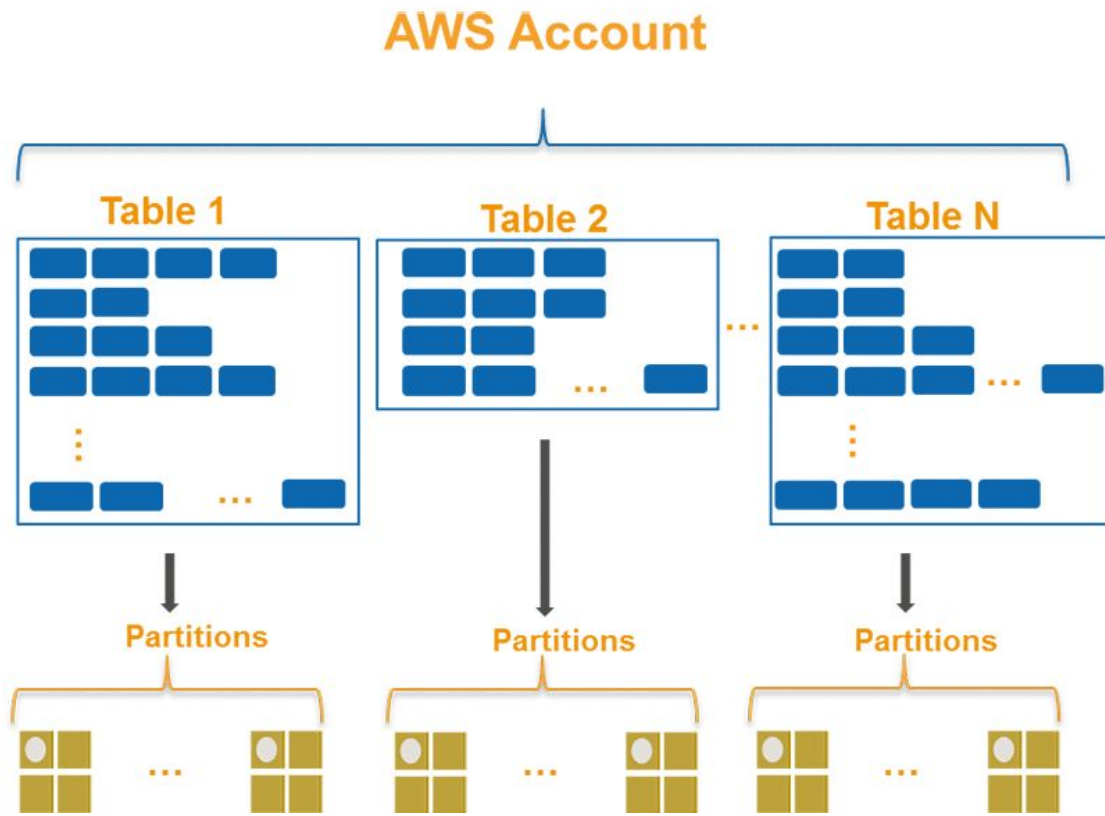
- **Recommendations** for partition keys:
 - Use high-cardinality attributes - **distinct** values for each item
 - Use composite attributes



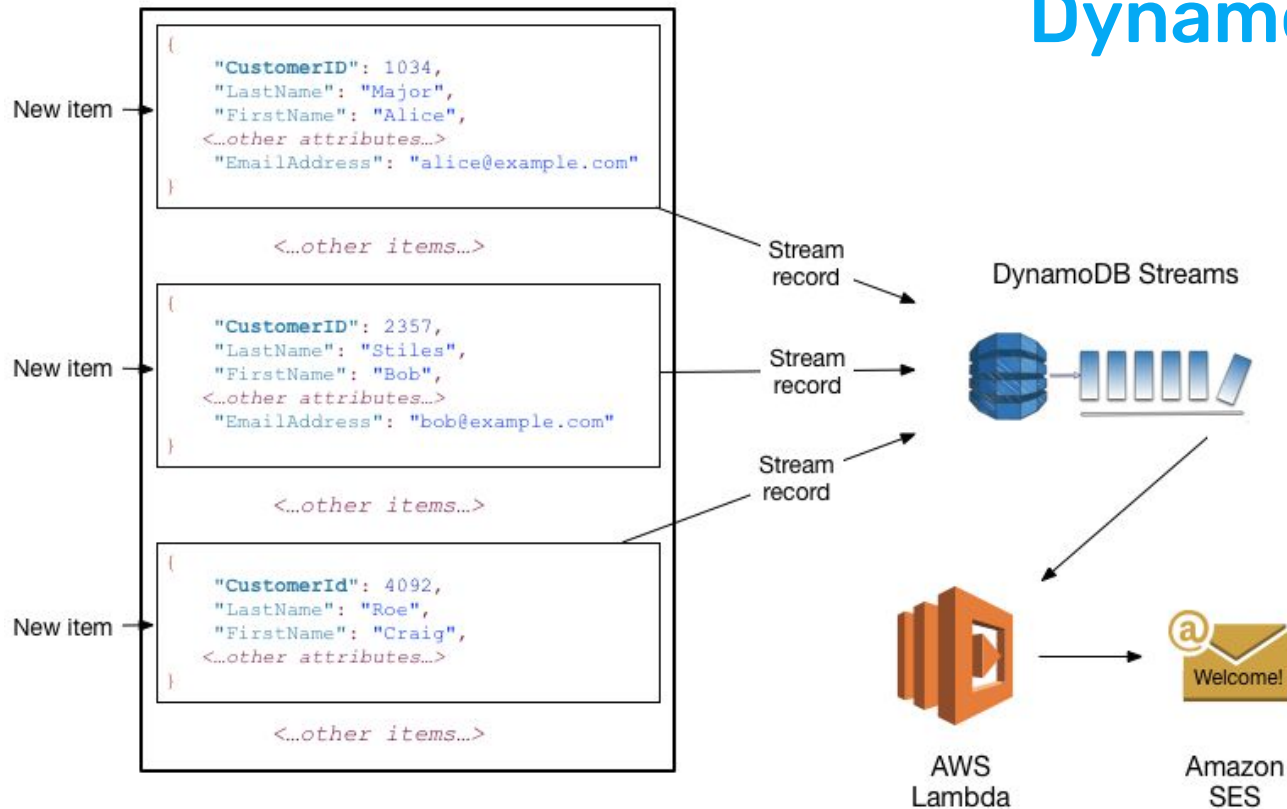
Partitioning

- Items are **distributed** across 10-GB storage units, called **partitions** (physical storage internal to DynamoDB)
- Using low-cardinality attributes as the partition key and order_date as the sort key greatly increases the likelihood of **hot partition issues**.

For example, if one product is more popular, then the reads and writes for that key is high, resulting in throttling issues.



Customers



DynamoDB Streams

Query vs Scan

Scan	Query
No need to specify any key criteria	Need to specify Partition Key mandatorily
Navigates through all the items in a table	Navigates through all the items in a partition
Maximum limit of 1 MB per page scanned	Maximum limit of 1 MB per page queried
FilterExpression operation can be used to narrow down the results, post scan	Sort key can be specified to narrow down the results of query. In addition, FilterExpression operation can also be used to narrow down the results, post query



Avoid full table scans!

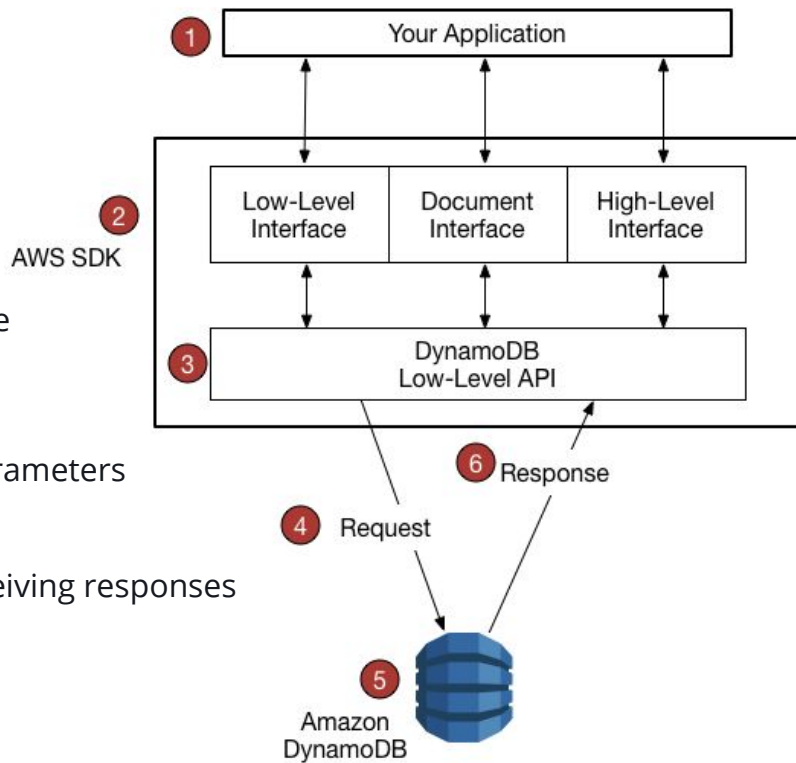
Other core concepts

- **Secondary indexes** [guide](#), [explained](#)
- **Strong** vs **Eventual** consistency
- Read/Write Capacity Mode [guide](#)


	Eventual Consistency	Strong consistency
Consistency	Propagation of latest update might take a few ms longer. It is possible to miss the latest update	You always read the latest update
Performance	Fastest possible reads	Slower than eventually consistent reads
Cost	Cheapest possible reads. Two eventually consistent reads cost 1 RCU	Twice as expensive as eventually consistent reads. Each strongly consistent read cost 1 RCU

AWS SDK for DynamoDB

- How it works?
 - The AWS SDK constructs/sends HTTP(S) requests for use with the low-level DynamoDB API
 - DynamoDB executes the request. Returns an HTTP code
- AWS SDKs provides important services:
 - Formatting HTTP(S) requests and serializing request parameters
 - Generating a cryptographic signature for each request
 - Forwarding requests to a DynamoDB endpoint and receiving responses from DynamoDB.
 - Extracting the results from those responses
 - Implementing basic retry logic in case of errors



Getting Started




- <https://www.dynamodbguide.com/> - Learn DynamoDB
- [Setting Up DynamoDB](#) - To learn how to set up DynamoDB (the downloadable version or the web service).
- [Working with DynamoDB](#) - Tables, Items, Queries, Scans, and Indexes
- JavaScript SDK [Documentation](#) & [Examples](#)

DynamoDB DEMO

AWS Cognito

Overview

- Secure and scalable user directory
 - Social and enterprise identity federation
 - Standards-based authentication
 - Security for your apps and users
 - Access control for AWS resources
 - Easy integration with your app
 - Can be used as a standalone IdP
-



User Pools

- Sign-up and sign-in services, social sign in
- A built-in, customizable web UI
- Forgot password flow, email or phone number verification, MFA
- User directory management and user profiles
- Security features
- Customized workflows and user migration through AWS Lambda triggers.

General settings

Users and groups
Attributes
Policies
MFA and verifications
Advanced security
Message customizations
Tags
Devices
App clients
Triggers
Analytics

App integration

App client settings
Domain name
UI customization
Resource servers

Federation

Identity providers
Attribute mapping

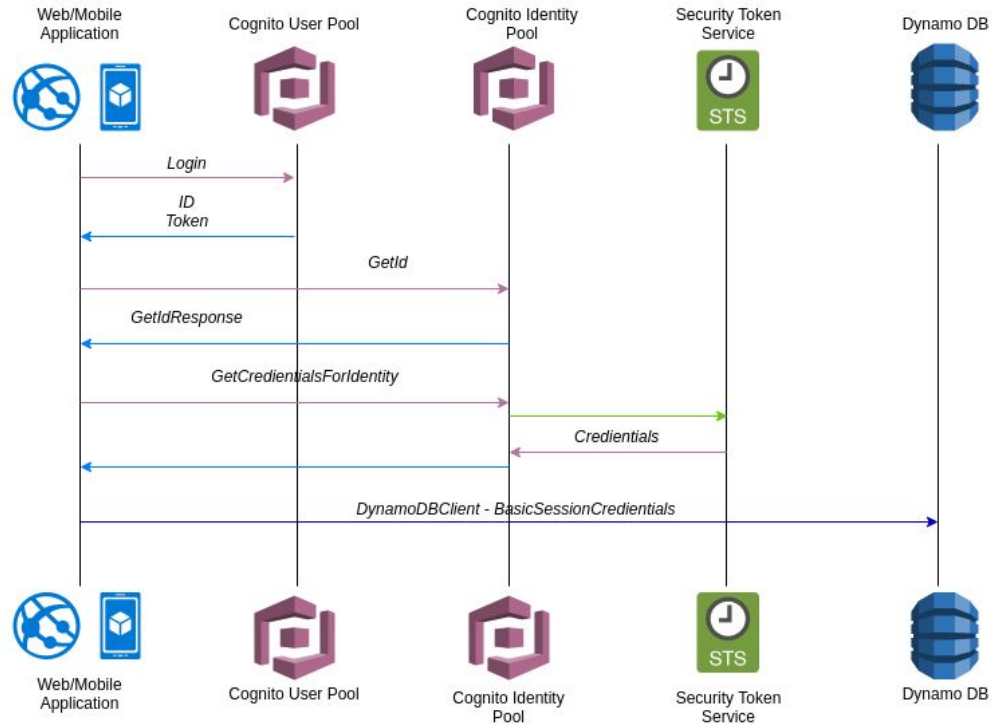
Admin Capabilities

- Create and manage User pools
- Define custom attributes
- Require Submission of Attribute Data
- Set per-app permissions
- Set up Password Policies
- Search users
- Manage users

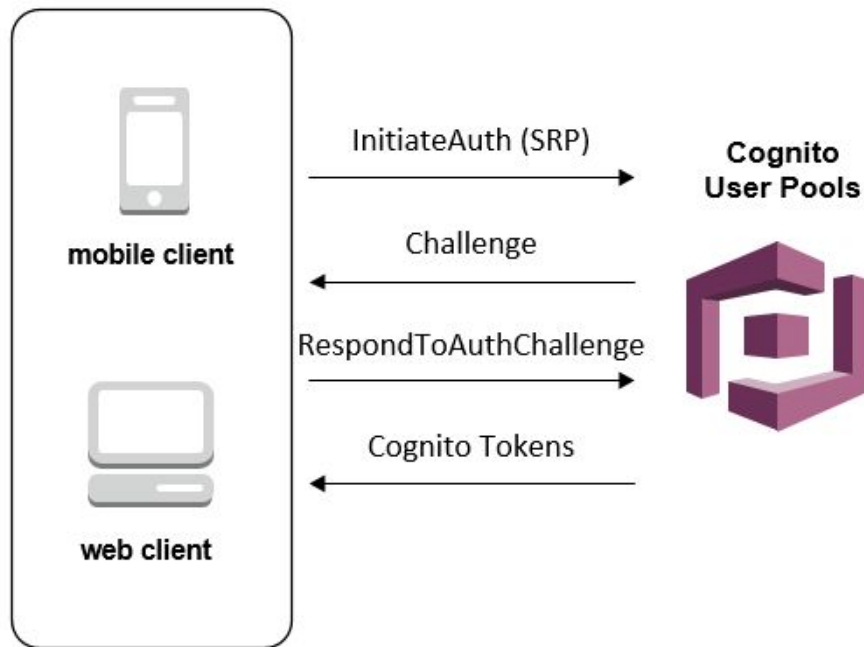


Federated Identities

- Public Providers (Amazon, Google, Facebook)
- Amazon Cognito User Pools
- Open ID Connect & SAML Identity
- Developer Authenticated Identities



Simple Authentication



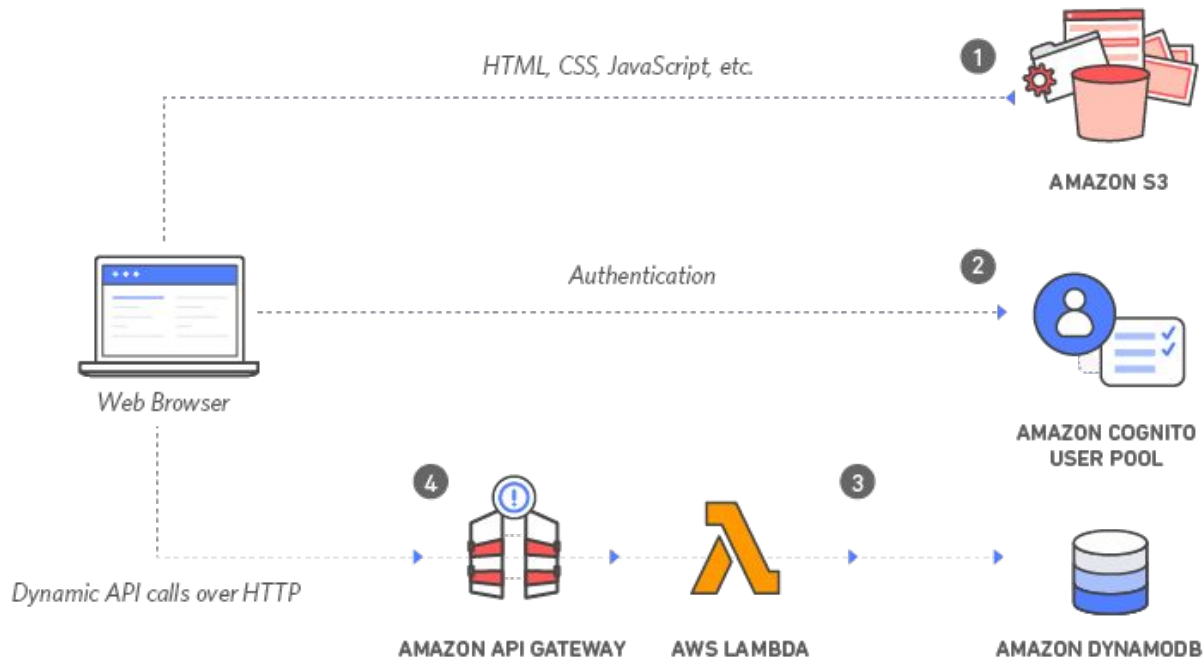
- [Getting Started](#) with User Pools
- SRP [Explained](#)
- [Understanding](#) Cognito for Serverless
- NPM [Module](#) for Node.js

Cognito Demo

Integrate **Cognito** with **API Gateway**

Wire up

- Wild Rydes [Tutorial](#)



Homework

1. DynamoDB Practice

Using AWS-SDK & JavaScript, create a DynamoDB Table for hobbies/skills. You can use a shared table with another Trainee ;) Decide what your partition/sort keys should look-like.

Example schema:

Name - *String*

Description - *String*

Practitioner - Your name

Since - *Date*

Rating - *Integer from 1 to 10*

Expertise - *ENUM ('novice', 'advanced beginner', 'competent', 'proficient', 'expert')*

Create a simple REST API for managing these hobbies (CRUD operations).

Create endpoints that return data via more advanced queries, for example:

- Show practitioner with the most hobbies
- Get top skill per practitioner (user)
- List top 5 favorite hobbies / activities / skills
- ... be creative!

